

Methods wait and notifyAll

We provide an example that shows the need for two more thread methods: `wait()` and `notifyAll()`. We implement a “dropbox”, a place where threads, called *producers*, can deposit an integer for other threads, called *consumers*, to pick up. A dropbox is like a vending machine that can hold only one item, *one* candy bar, and if there is no candy bar in it, you have to wait until the vendor puts one there.

A producer wanting to put a new integer into a full dropbox must wait until a consumer takes the current integer out. A consumer wanting to take out an integer must wait until there is one there.

To implement waiting, Java maintains *two* lists for a synchronized object *sob*.

1. The *locklist*: The list of threads that are waiting to obtain the lock on *sob*; they are ready to execute a synchronized statement or method.
2. The *waitlist*: The list of threads that *had* the lock but couldn't proceed because *sob* was not in a suitable state (e.g. a consumer found the dropbox empty). So they executed method `wait()`. They are waiting until the state of *sob* changes.

Suppose a thread *t* uses a synchronized statement or method to acquire the lock on *sob*. Two things can happen:

1. Thread *t* can't complete its task (e.g. a consumer finds the dropbox empty). Therefore, *t* executes a call `wait()`, which puts *t* on the *waitlist* and makes it relinquish the lock on *sob*. Thus, another thread can obtain the lock. Of course, when *t* is given another chance, *t*'s call on `wait()` completes and *t* continues executing from that point.
2. Thread *t* can complete its task. It does so and executes a call `notifyAll()`. This call moves all threads from the *waitlist* to the *locklist*. Why? Thread *t* changed the state of *sob*, so threads on the *waitlist* may now be able to complete their task. Moving them to the *locklist* gives them a chance to try again.

Class Dropbox

Class `Dropbox` has two fields: `full` tells whether the `Dropbox` contains an integer; if it does, the integer is in field `box`.

The two calls on `wait()` are in try-statements; this is necessary because an `InterruptedException` may be thrown while the method is waiting.

A consumer will call method `take()`. Suppose the call gets the lock and the `Dropbox` is full. The while-loop terminates. Field `full` is set to indicate that the `Dropbox` is empty, `notifyAll()` is called to move all threads on the *waitlist* to the *locklist*, and the value in `box` is returned.

Suppose the call `take()` gets the lock but the `Dropbox` is empty. Then `wait()` is called, so the thread is put on the *waitlist* and gives up the lock. When it is given another chance, execution `wait()`, the try-statement, and the repeatend all terminate and the while-condition `!full` is evaluated again.

Why is a while-loop necessary? Suppose an if-statement is used instead; the method body is:

```
if (!full)
    try {wait();} catch (...) {}
full= false; notifyAll(); return box;
```

Consider this execution sequence: (1) `take()` is called and the dropbox is empty. The call `wait()` is executed. (2) A producer puts an integer into the dropbox and does `notifyAll()`. (3) A *different* consumer is given the lock and takes the value out, thus emptying the buffer. (4) This consumer is given the lock; its call `wait()` completes, and it continues as if the dropbox were full, but it is not. The while loop is needed to make sure that a consumer processes normally *only* if the dropbox is full.

Another reason for the while-loop: If an `InterruptedException` is thrown and caught, `full` might be false.

```
/** An instance is a dropbox: A place for
 * producers to store an integer and for
 * consumers to take it out. */
class Dropbox {
    private boolean full= false; // = "box is full"
    private int box; // dropbox value (if full)

    /** Wait for box to hold an integer;
     * then take it out and return it. */
    public synchronized int take() {
        while (!full) {
            try {wait();}
            catch (InterruptedException ex) {}
        }
        full= false; notifyAll(); return box;
    }

    /** Wait for box to be empty; put n into it. */
    public synchronized void put(int n) {
        while (full) {
            try {wait();}
            catch (InterruptedException e) {}
        }
        box= n; full= true; notifyAll();
    }
}
```

Methods wait and notifyAll

A producer calls method `put(int)` to put an integer into the `Dropbox`, but if the `Dropbox` is full, the producer must wait until it is empty. The structure of method `put(int)` is similar to that of method `take()`, and we leave you to study it.

A complete application

To the right, we show a class `Main` that creates two `Consumer` threads and one `Producer` thread and starts all three executing.

Class `Consumer` alternately takes a value from a `Dropbox` and sleeps. Class `Producer` is similar.

The producers and consumers are given the `Dropbox` with which to work as a parameter of their constructors. This is a typical way to write such classes. It is more flexible than the method used in pdf file 6 under the `JavaHyperText` entry for `Threads`, which used instead a public static variable declared in the main class.

We have placed a `println` statement in the while-loop of the `Consumer`'s method `run()`. Put these classes into an Eclipse or DrJava project and execute this application and see what happens. Add other `println` statements to class `Dropbox`, if you want, in order to see how the `Producer` and two `Consumers` work together.

What about method `notify()`?

Method `notifyAll()` moves *all* threads from the *waitlist* to the *locklist*. There is a method `notify()`, which instead moves only *one* thread, chosen arbitrarily, from the *waitlist* to the *locklist*. Using `notify()` instead of `notifyAll()` may work, but in some cases it causes deadlock — no process can make progress. Don't use `notify()` unless you know what you are doing.

The following document in `JavaHyperText` entry *Threads* explains why `notify()` may not work:

Warning: use `notifyAll` and not `notify` unless you know what you are doing.

```
import java.util.Random;

public class Main {
    /** Create two Consumer threads and a
     * Producer thread and start all three. */
    public static void main(String[] args) {
        Dropbox box= new Dropbox();
        new Thread(new Consumer(box)).start();
        new Thread(new Consumer(box)).start();
        new Thread(new Producer(box)).start();
    }
}

/** An instance alternately takes an integer
 * from a Dropbox and sleeps */
public class Consumer implements Runnable {
    private Dropbox box;

    /** Constructor: a Consumer using db */
    public Consumer(Dropbox db) {box= db; }

    /** Forever: Get a value from the Dropbox
     * and sleep for a random time. */
    public void run() {
        Random random= new Random();
        while (true) {
            int i= box.take();
            System.out.println(
                Thread.currentThread().getName() +
                " " + i);
            try {
                Thread.sleep(random.nextInt(100));
            } catch (InterruptedException e) { }
        }
    }
}

/** An instance repeatedly sleeps and
 * puts a random number into a Dropbox. */
class Producer implements Runnable {
    private Dropbox box;

    /** Constructor: a Producer using db */
    public Producer(Dropbox db) {box= db;}

    /** Forever: sleep for a random time and
     * put a random number into the Dropbox. */
    public void run() {
        Random random= new Random();
        while (true) {
            int n= random.nextInt(10);
            try {
                Thread.sleep(random.nextInt(100));
            } catch (InterruptedException e) { }
            box.put(n);
        }
    }
}
```